

## Lecture 4: Sequencing Coverage Analysis

### + Introduction to Genome-Wide Association Studies (GWAS)



ECE 365 - Data Science and Genomics

## Announcements:

- Lab 1 due today (11:59pm)
- Lab 2 released (due March 18)


## Today:

- Finish discussion on sequencing coverage
- Introduction to Genome-Wide Association Studies (GWAS)

# Indexing

16 read  
cgtcagcggacagggc

ggtttaatgtggttctgcttggcggtagtcattaaagagccccgtggtggccaat  
caagaaaatgtcacgcgcgttcccagcactttcagctggtttgtcgtagcccat  
caccaccgtaagccaagaccagcttcaggccaagtagccttccgccagcgggt  
ctgcgtcggcatggattctgcacggcaaaagttcacgcgtcggtttgccataatt  
aaggacgcgcctggattcaccttgcgatcggcaatcgcaggaatgagagagcag  
ataatgaaagcgttgacgtaagaaaagccatcggtttcccggtaccgggttttgc  
gctgccccggtcagtcagcgacctcgccagcgtcagcggacagggcgcaagt  
ccgtgaatgggccgtacagttatgaaacccttttttctaaagggccttctacaa  
cccttggtatgcagggcgaagtcgggaaaacttctgttctgtttaaaatgtggtt  
tgctcatagtgtggtagatctcagcttactattggctttaacgaaagccgtatt  
ccggtgaaaaataacagtcacgcgttttagttgttaatgttacaccaacaacgaaa  
ccaacacgccaggcttaattcctgtggagttatatatgagcgtaaatcggatcc



all substrings  
of length 16

Dict()

ggtttaatgtggttct	-->	0
gtttaatgtggttctg	-->	1
tttaatgtggttctgc	-->	2
ttaatgtggttctgct	-->	3
taatgtggttctgctt	-->	4
aatgtggttctgcttg	-->	5
atgtggttctgcttgg	-->	6
tgtggttctgcttggc	-->	7
gtggttctgcttggcg	-->	8
tgggttctgcttggcgg	-->	9
gggttctgcttggcggt	-->	10
gttctgcttggcggtta	-->	11
ttctgcttggcggttag	-->	12
tctgcttggcggttagt	-->	13
ctgcttggcggttagtc	-->	14
tgcttggcggttagtca	-->	15
gcttggcggttagtcat	-->	16
cttggcggttagtcatt	-->	17
ttggcggttagtcatta	-->	18

- Another way to do indexing: Python dictionary (hash function)
- Let's look at this in a notebook

# What if there are errors/mutations on read?

$X =$  

$Y =$

ggtttaatgtggttctgcttgccggtagtcattaagagccccgtggtggccaat  
caagaaaatgtcacgcgcgttcccagcactttcagctgttttgtcgtagcccat  
caccaccgtaagccaagaccagcttcaggccaagtagccttccgccagcgggt  
ctgcgtcggcatggattctgcacggcaaagttcacgcgtcggtttgccataatt  
aaggacgcgcctggattcacottgcgatcggcaatcgcaggaatgagagagcag  
ataatgaaagcgttgacgtaagaaagccatcgttttcccggtaccggtttttgc  
gcctgcccggtacgtcagcgacctcgccagcgtcagcggacaggcgcaagtgc  
ccgtgaatgggccgtacagttatgaaacccttttttctaaagggccttctacaa  
ccottggatgcagggcgaagtcgggaaaacttctgttctgtttaaaatgtgttt  
tgctcatagtgtggtagatctcagcttactattggctttaacgaaagccgtatt  
ccggtgaaaataacagtcacgcttttagttgttaatgttacaccaacaacgaaa  
ccaacacgccaggcttaattcctgtggagttatatatgagcgtaaatacgatcc

# What if there are errors/mutations on read?

$X = \text{cgtaagcggacatggc}$

$Y =$


```
ggtttaatgtggttctgcttgcggtagtcattaagagccccgtggtggccaat
caagaaaatgtcacgcgcgttcccagcactttcagctgttttgtcgtagcccat
caccaccgtaagccaagaccagcttcaggccaagtagccttccgccagcgggt
ctgcgtcggcatggattctgcacggcaaagttcacgcgtcggtttgccataatt
aaggacgcgcctggattcacottgcgatcggcaatcgcaggaatgagagagcag
ataatgaaagcgttgacgtaagaaagccatcgttttcccggtaccggttttgc
gcctgccgggtacgtcagcgacctcgccagcgtcagcggacagggcgcaagtg
ccgtgaatgggccgtacagttatgaaacccttttttctaaggggcttctacaa
ccottggatgcagggcgaagtcgggaaaacttctgttctgtttaaaatgtgttt
tgctcatagtgtggtagatctcagcttactattggctttaacgaaagccgtatt
ccggtgaaaataacagtcacgcgttttagttgtaatgttacaccaacaacgaaa
ccaacacgccaggcttaattcctgtggagttatatagagcgtaaatcgatcc
```

- One idea: Consider indexing substrings of length  $k \approx 6$

# What if there are errors/mutations on read?

$X =$   `cgtaagcggacatggc`

$Y =$  `ggtttaatgtggttctgcttggcggtagtcattaagagccccgtggtggccaat  
caagaaaatgtcacgcgcgttcccagcactttcagctgttttgtcgtagcccat  
caccaccgtaagccaagaccagcttcaggccaagtagccttccgccagcggtt  
ctgcgtcggcatggattctgcacggcaaaagttcacgcgtcggtttgccataatt  
aaggacgcgcctggattcaccttgcgatcggcaatcgaggaatgagagagcag  
ataatgaaagcgttgacgtaagaaagccatcgttttcccggtaccggttttgc  
gcctgccgggtacgtcagcgacctcgccagcgtcagcggacaggcgcaagtgc  
ccgtgaatgggccgtacagttatgaaaccttttttctaaagggccttctacaa  
cccttgatgcagggcgaagtcgggaaaaacttctgtctgtttaaaatgtgttt  
tgctcatagtgtggtagatctcagcttactattggctttaacgaaagccgtatt  
ccggtgaaaataacagtcacgcttttagttgtaatgttacaccaacaacgaaa  
ccaacacgccaggcttaattcctgtggagttatatatgagcgtaaatcgatcc`



all substrings  
of length  $k = 6$

Python Dict ( )

```
ggttta --> [0]
gtttaa --> [1, 471]
tttaat --> [2]
ttaatg --> [3, 571]
taatgt --> [4, 572]
aatgtg --> [5, 477]
atgtgg --> [6]
tgtggg --> [7, 495]
gtgggt --> [8]
tgggtc --> [9]
ggttct --> [10, 158]
gttctg --> [11, 159, 466]
ttctgc --> [12, 160, 177]
tctgct --> [13]
ctgctt --> [14]
tgcttg --> [15]
gcttgg --> [16]
cttggc --> [17]
ttggcg --> [18]
tggcgg --> [19]
ggcggg --> [20]
```

- One idea: Consider indexing substrings of length  $k$

# What if there are errors/mutations on read?

$X =$  **cg****t****a****agcggacat****g****gc**

take substrings  
of length  $k = 6$

cgtaag --> [114] 286]  
gtaagc --> [115]  
taagcg -->  
aagcgg -->  
agcgga --> [359]  
gcggac --> [360]  
cggaca --> [361]  
ggacat -->  
gacatg -->  
acatgg -->  
catggc -->

## What if there are errors/mutations on read?

$X =$  cgt**a**gcggacat**t**ggc

take substrings  
of length  $k = 6$

cgtaag --> [114] 286]  
gtaagc --> [115]  
taagcg -->  
aagcgg -->  
agcgga --> [359]  
gcgga --> [360]  
cgga --> [361]  
ggacat -->  
gacatg -->  
acatgg -->  
catggc -->

Smith Waterman  $(X, Y)$

genome

→ bad alignment

...tagcccatcaccaccgtaagcgaagacccagcttcaggccaagtagccttcgccagcgg...  
 ↑  
 114

Handwritten notes: "cgt aag cgg acat ggc" with arrows pointing to "5' alignment" and "3' alignment".

© 2011. Waterman

smith Waterman (X, Y)  
good alignment!

...ccggctacgtcagcgacctcgccagcgtcagcggacagggcgcaagtgccgtgaatgggc...




# How do we choose $k$ ?

$$\lceil 3.2 \rceil = 4$$

$$\lceil 5 \rceil = 5$$

- Based on the sequencing technology (e.g., Illumina:  $L = 100$ , error rate = 0.1%)

. suppose at most 1 error per read



Claim: there is a segment of length  $\frac{100-1}{2} = 49.5$  with no errors

. suppose  $\leq t$  errors per read



there is a segment of length

$$\left\lceil \frac{L-t}{t+1} \right\rceil \text{ with no errors}$$

$\hookrightarrow \text{set } = k$

How many reads have  $> t$  errors?

# of errors in a read:  $X \sim \text{Binomial}(100, 0.001)$

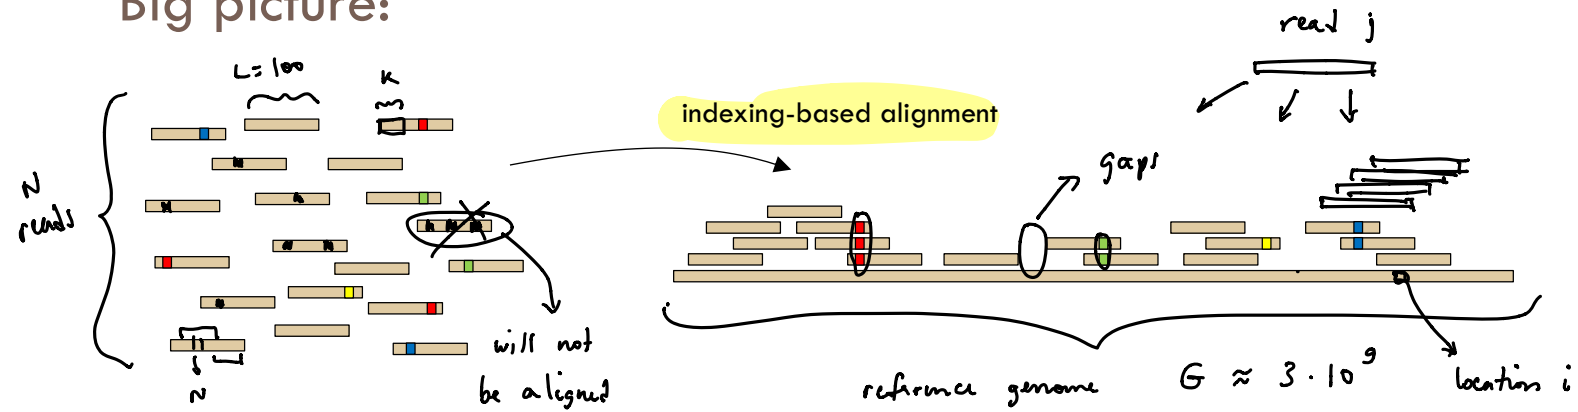
fraction of reads with  $> t$  errors:  $P(X > t)$

e.g. for  $t = 2$ .

$$P(X > 2) = 0.00015$$

$$k = \left\lceil \frac{100-2}{3} \right\rceil = 33$$

# Big picture:




How to pick  $N$ ?

Assume that start location of each read is uniformly distributed on genome

$$P(\text{read } j \text{ covers position } i) = \frac{L}{G}$$

$$\Rightarrow \text{Expected \# of reads covering position } i = \frac{NL}{G} \triangleq c \quad \left( \begin{array}{l} \text{coverage} \\ \text{depth} \end{array} \right)$$

Prob. that there are gaps? Fix  $i$ .



$$P(i \text{ is uncovered}) = \prod_{j=1}^N \left(1 - \frac{L}{G}\right) = \left(1 - \frac{L}{G}\right)^N = \underbrace{\left(1 - \frac{1}{\frac{G}{L}}\right)^{\frac{G}{L}}}_{\approx e^{-1}} \cdot \underbrace{\frac{NL}{G}}_c \approx e^{-c}$$

$$\text{Recall: } \lim_{n \rightarrow \infty} \left(1 - \frac{1}{n}\right)^n = e^{-1}$$

union bound (good approximation)


$$P(\text{some position is uncovered}) \stackrel{\downarrow}{\leq} G e^{-c} = \epsilon \quad (\text{desired failure prob.})$$

$$\Rightarrow e^c = \frac{G}{\epsilon} \Rightarrow c = \ln\left(\frac{G}{\epsilon}\right)$$

$$\text{Since } c = \frac{NL}{G}, \quad N = \frac{G}{L} \ln\left(\frac{G}{\epsilon}\right). \quad (\text{Lander - Watterman})$$


# Big picture:

- Map sequencing data to a reference genome



# Big picture:


- Map sequencing data to a reference genome



- Using aligned data, we can identify *variants*

# Big picture:

- Map sequencing data to a reference genome



- Using aligned data, we can identify *variants*
- Variants allow us to create genotype matrices


different variants

person 1	0	1	1	0	0	1
	1	0	0	1	0	0
	1	1	1	0	0	0
$\vdots$	0	0	0	0	1	1
	1	1	1	0	0	0
	0	0	0	0	1	0
person $n$	1	1	0	1	0	1

genotype


# Genotype data

- Focus on a set of common variants in the genome



# Genotype data

- Focus on a set of common variants in the genome




- Typically, we focus on Single-Nucleotide Polymorphisms (SNPs, read snips)



# Genotype data


- Focus on a set of common variants in the genome



- Typically, we focus on Single-Nucleotide Polymorphisms (SNPs, read snips)
- Most SNPs only allow two possible values: REF and ALT (e.g. REF = A, ALT = C)

# Genotype data

- Focus on a set of common variants in the genome



- Typically, we focus on Single-Nucleotide Polymorphisms (SNPs, read snips)
- Most SNPs only allow two possible values: REF and ALT
- More info on SNPs: SNPedia!
  - Search for rs429358

**SNPedia**

Talk ? Contributions Create account Log in

Search SNPedia



Navigation ▾

Page Discussion Edit with form Edit History

Have questions? Visit <https://www.reddit.com/r/SNPedia>



**Empower yourself**

Order now

**rs429358**

This SNP, located in the fourth exon of the *ApoE* gene, affects the amino acid at position 130 of the resulting protein. The more common **rs429358** allele is (T). If the allele is (C) and the same chromosome also harbors the **rs7412**(C) allele, the combination is known as an **APOE-ε4** allele. The APOE-ε4 allele has a strong influence on the risk of **Alzheimer's disease**.

Many studies have estimated the level of risk, and it varies depending on age, sex, ethnicity,

Orientation plus  
Stabilized plus

Geno • Mag • Summary •

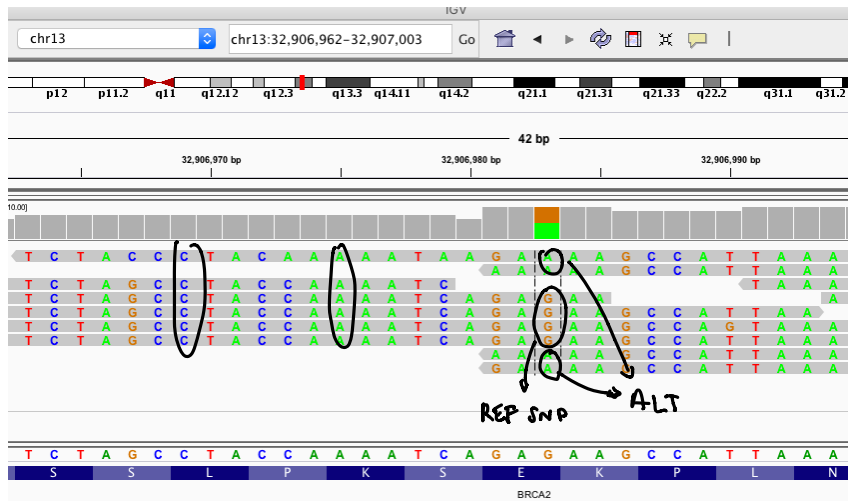
(C;C)

1.2

one of 2 snps relevant to  
classifying APOE genotype

Don't forget: humans are diploid → two copies of each chromosome

- Two “copies” of each chromosome (different variants)
- Aligned data is a combination from both:



# Variant data: VCF files

- Standard file format to list variants of one or many individuals
- Let's take a look!

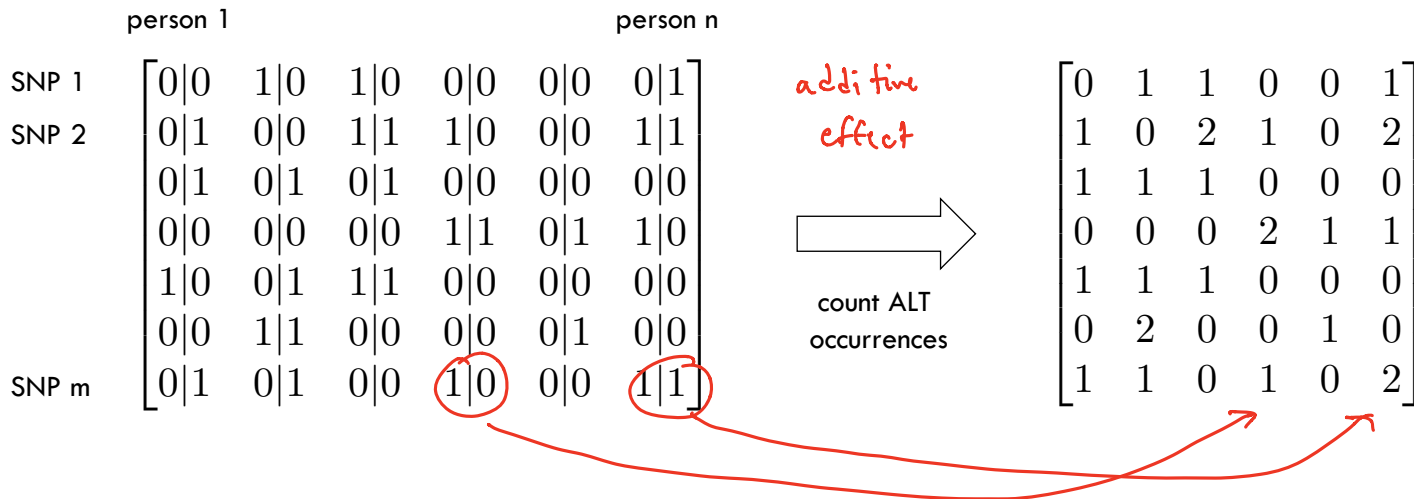
# Variant data: VCF files

- Standard file format to list variants of one or many individuals
- Let's take a look!
- Genotype matrix:

	person 1				person n	
SNP 1	0 0	1 0	1 0	0 0	0 0	0 1
SNP 2	0 1	0 0	1 1	1 0	0 0	1 1
	0 1	0 1	0 1	0 0	0 0	0 0
	0 0	0 0	0 0	1 1	0 1	1 0
	1 0	0 1	1 1	0 0	0 0	0 0
	0 0	1 1	0 0	0 0	0 1	0 0
SNP m	0 1	0 1	0 0	1 0	0 0	1 1

# Variant data: VCF files

- Standard file format to list variants of one or many individuals
- Let's take a look!
- Genotype matrix:



# Genome-Wide Association Studies (GWAS)

	SNP 1				SNP m	
person 1	0	1	1	0	0	1
	1	0	2	1	0	2
	1	1	1	0	0	0
	0	0	0	2	1	1
	1	1	1	0	0	0
	0	2	0	0	1	0
person n	1	1	0	1	0	2

# Genome-Wide Association Studies (GWAS)

	SNP 1				SNP m	
person 1	0	1	1	0	0	1
	1	0	2	1	0	2
	1	1	1	0	0	0
	0	0	0	2	1	1
	1	1	1	0	0	0
	0	2	0	0	1	0
person n	1	1	0	1	0	2

genome-wide

phenotype

$\begin{bmatrix} 0 \\ 0 \\ 1 \\ 1 \\ 1 \\ 0 \\ 1 \end{bmatrix}$

e.g. has Diabetes

(doesn't need to be  
binary. e.g., height)



# Genome-Wide Association Studies (GWAS)

*associated*

	SNP 1		SNP 100		SNP m		phenotype
person 1	0	1	1	0	0	1	0
	1	0	2	1	0	2	0
	1	1	1	0	0	0	1
	0	0	0	2	1	1	1
	1	1	1	0	0	0	1
	0	2	0	0	1	0	0
person n	1	1	0	1	0	2	1
<i>new person</i>	<i>1</i>	<i>0</i>	<i>1</i>	<i>1</i>	<i>2</i>	<i>0</i>	<i>?</i> (predict)

- Which SNPs are associated with a given phenotype?
- Given a new individual's genotype, can you predict their phenotype?

# Revisiting Logistic Regression

- Predict binary variable from real-valued features

$$p(1 | \underline{x}) = \frac{e^{\beta_0 + \underline{\beta}^T \underline{x}}}{1 + e^{\beta_0 + \underline{\beta}^T \underline{x}}} = \frac{1}{1 + e^{-(\beta_0 + \underline{\beta}^T \underline{x})}}$$

"   
 p

$$1 + e^{-(\beta_0 + \underline{\beta}^T \underline{x})} = \frac{1}{p} \Rightarrow e^{-(\beta_0 + \underline{\beta}^T \underline{x})} = \frac{1}{p} - 1 = \frac{1-p}{p}$$

$$\Rightarrow \ln\left(\frac{p}{1-p}\right) = \underbrace{\beta_0 + \underline{\beta}^T \underline{x}}_{\text{linear model}}$$

$\frac{p}{1-p}$  : odds ratio

$\ln \frac{p}{1-p}$  : log odds ratio