

Lecture 5: GWAS (cont.)




ECE 365 - Data Science and Genomics

Announcements:

- Lab 2 (Sequence alignment) due on Thursday
- Lab 3 (GWAS) released tomorrow


Genotype data

- Focus on a set of common variants in the genome



Genotype data

- Focus on a set of common variants in the genome




- Genotype data (VCF file) looks like this:

person 1 *paternal* *maternal* person n

SNP 1	0 0	1 0	1 0	0 0	0 0	0 1
SNP 2	0 1	0 0	1 1	1 0	0 0	1 1
	0 1	0 1	0 1	0 0	0 0	0 0
	0 0	0 0	0 0	1 1	0 1	1 0
	1 0	0 1	1 1	0 0	0 0	0 0
	0 0	1 1	0 0	0 0	0 1	0 0
SNP m	0 1	0 1	0 0	1 0	0 0	1 1

Genotype data

- Focus on a set of common variants in the genome



- Genotype data (VCF file) looks like this:

	person 1						person n					
SNP 1	0 0	1 0	1 0	0 0	0 0	0 1	0	1	1	0	0	1
SNP 2	0 1	0 0	1 1	1 0	0 0	1 1	1	0	2	1	0	2
	0 1	0 1	0 1	0 0	0 0	0 0	1	1	1	0	0	0
	0 0	0 0	0 0	1 1	0 1	1 0	0	0	0	2	1	1
	1 0	0 1	1 1	0 0	0 0	0 0	1	1	1	0	0	0
	0 0	1 1	0 0	0 0	0 1	0 0	0	2	0	0	1	0
SNP m	0 1	0 1	0 0	1 0	0 0	1 1	1	1	0	1	0	2

count ALT occurrences

Genome-Wide Association Studies (GWAS)

	SNP 1				SNP m	
person 1	0	1	1	0	0	1
	1	0	2	1	0	2
	1	1	1	0	0	0
	0	0	0	2	1	1
	1	1	1	0	0	0
	0	2	0	0	1	0
person n	1	1	0	1	0	2

genome-wide

phenotype

$\begin{bmatrix} 0 \\ 0 \\ 1 \\ 1 \\ 1 \\ 0 \\ 1 \end{bmatrix}$

e.g. has Diabetes

(doesn't need to be
binary. e.g., height)

Genome-Wide Association Studies (GWAS)

	SNP 1		SNP ₁₀₀		SNP m		phenotype
person 1	0	1	1	0	0	1	0
	1	0	2	1	0	2	0
	1	1	1	0	0	0	1
	0	0	0	2	1	1	1
	1	1	1	0	0	0	1
	0	2	0	0	1	0	0
person n	1	1	0	1	0	2	1
new person	1	0	1	1	2	0	? (predict)

- Which SNPs are associated with a given phenotype?
- Given a new individual's genotype, can you predict their phenotype?

Revisiting Logistic Regression

- Predict binary variable from real-valued features

$$p(1 | \underline{x}) = \frac{e^{\beta_0 + \beta^T \underline{x}}}{1 + e^{\beta_0 + \beta^T \underline{x}}} = \frac{1}{1 + e^{-(\beta_0 + \beta^T \underline{x})}}$$

$$1 + e^{-(\beta_0 + \beta^T \underline{x})} = \frac{1}{p} \Rightarrow e^{-(\beta_0 + \beta^T \underline{x})} = \frac{1}{p} - 1 = \frac{1-p}{p}$$

$$\Rightarrow \ln\left(\frac{p}{1-p}\right) = \underbrace{\beta_0 + \beta^T \underline{x}}_{\text{linear model}}$$

$$\frac{p}{1-p} : \text{odds ratio} \in (0, \infty)$$

$$\ln \frac{p}{1-p} : \log \text{ odds ratio} \in (-\infty, \infty)$$

Revisiting Logistic Regression

- Can we use Logistic Regression for GWAS?

	mills						
	SNP 1			SNP m			phenotype
person 1	0	1	1	0	0	1	$\begin{bmatrix} 0 \\ 0 \\ 1 \\ 1 \\ 1 \\ 0 \\ 1 \end{bmatrix}$
	1	0	2	1	0	2	
x_i	1	1	1	0	0	0	
	0	0	0	2	1	1	
	1	1	1	0	0	0	
	0	2	0	0	1	0	
person n	1	1	0	1	0	2	

$$p(1 | \underline{x}) = \frac{1}{1 + e^{-(\beta_0 + \beta^T \underline{x})}}$$

↙
risk of having a disease

Revisiting Logistic Regression

- Can we use Logistic Regression for GWAS?

millions

	SNP 1					SNP m	phenotype
person 1	0	1	1	0	0	1	0
	1	0	2	1	0	2	0
	1	1	1	0	0	0	1
	0	0	0	2	1	1	1
	1	1	1	0	0	0	1
	0	2	0	0	1	0	0
person n	1	1	0	1	0	2	1

1000s

- Problem: Number of SNPs can be $\sim 10^6$

Revisiting Logistic Regression

- Can we use Logistic Regression for GWAS?

	SNP 1				SNP m		phenotype
person 1	0	1	1	0	0	1	0
	1	0	2	1	0	2	0
	1	1	1	0	0	0	1
	0	0	0	2	1	1	1
	1	1	1	0	0	0	1
	0	2	0	0	1	0	0
person n	1	1	0	1	0	2	1

- Problem: Number of SNPs can be $\sim 10^6$

GWAS via *univariate* logistic regression

- Run a separate logistic regression for each SNP

	SNP 1			x_i	SNP m			phenotype
person 1	0	1	1	0	0	1		0
	1	0	2	1	0	2		0
	1	1	1	0	0	0		1
	0	0	0	2	1	1		1
	1	1	1	0	0	0		1
	0	2	0	0	1	0		0
person n	1	1	0	1	0	2		1

$$p(1 | x_i) = \frac{1}{1 + e^{-(\beta_0 + \beta_i x_i)}}$$

Captures association
between ith SNP
and phenotype

- Use this to identify small subset of SNPs associated with phenotype
- Let's look at some examples on a Jupyter notebook

GWAS via *univariate* logistic regression

- We will get a β for each SNP

	SNP 1						SNP m	phenotype
person 1	0	1	1	0	0	1		0
	1	0	2	1	0	2		0
	1	1	1	0	0	0		1
	0	0	0	2	1	1		1
	1	1	1	0	0	0		1
	0	2	0	0	1	0		0
person n	1	1	0	1	0	2		1
	β_1	β_2	...			β_m		

model for *ith* SNP:

$$\ln\left(\frac{p(1|x_i)}{1-p(1|x_i)}\right) = \beta_0 + \beta_i x_i$$

↓
LOR for reference
genome (no
ALT SNPs)

GWAS via *univariate* logistic regression

- Idea: combine all beta coefficients into a single model:

$$\ln\left(\frac{p(1|x)}{1 - p(1|x)}\right) = \beta_0 + \beta_1 x_1 + \cdots + \beta_m x_m$$

GWAS via *univariate* logistic regression

- Idea: combine all beta coefficients into a single model:

$$\ln\left(\frac{p(1|x)}{1 - p(1|x)}\right) = \beta_0 + \beta_1 x_1 + \cdots + \beta_m x_m$$

- Problems with this approach:

GWAS via *univariate* logistic regression

- Idea: combine all beta coefficients into a single model:

$$\ln\left(\frac{p(1|x)}{1 - p(1|x)}\right) = \beta_0 + \beta_1 x_1 + \dots + \beta_m x_m$$

- Problems with this approach:
 - ▣ Since m is very large, some β_i s will be large **by chance**

GWAS via *univariate* logistic regression

- Idea: combine all beta coefficients into a single model:

$$\ln\left(\frac{p(1|x)}{1-p(1|x)}\right) = \beta_0 + \beta_1 x_1 + \dots + \beta_m x_m$$

- Problems with this approach:
 - ▣ Since m is very large, some β_i s will be large **by chance**
 - ▣ Some x_i s are correlated

(e.g., anyone with $x_3 = 1$ has $x_4 = 1$)

Identifying statistically significant SNPs

- To measure the significance of the association, we use the p -value

Identifying statistically significant SNPs


- To measure the significance of the association, we use the *p-value*
 - Probability that the coefficient β_i would be obtained by chance if there was **no** association

$$P(|\hat{\beta}_i| > |\beta_i|_{\text{found}} \mid \text{no association})$$

- We can use the statsmodels Python package to perform the logistic regressions and compute the *p-values*
- Let's return to the jupyter notebook

Manhattan plots

- Allow us to see the significance of all SNPs in the genome
- We plot $-\log_{10}(p\text{-value})$



Manhattan plot of SNP association with Irritable Bowel Disease

