

# ECE365: Introduction to NLP

Spring 2021

Lecture 7: Word Embeddings

[Reading J&M Chapter 6]

# Logistics

- Quiz 2 Tuesday 05/04, 05/05

# Distributional Hypothesis

**Distributional hypothesis**, stated by linguist John R. Firth (1957) as:

“You shall know a word by the company it keeps.”

≈ “words that occur in similar contexts tend to have similar meanings”

One of the most successful ideas of modern statistical NLP!

# Distributional Semantic Models

- Vector semantics = {distributional idea (defining a word by counting what other words occur in its environment) } + {meaning of a word as a vector (a point in N-dimensions)}
- Popularly called **word embeddings**
  - Various versions depending on how the vector components are computed
  - Latent Semantic Analysis, Word2vec, GloVe

# Why model words as vectors?

- We need to model word meaning
- As a way for computing similarity between words
- Useful for unknown words

# Why model words as vectors?

- We need to model word meaning
- As a way for computing similarity between words
- Similar words
  - not synonyms, but share some element of meaning
  - Car- bicycle, cow-horse

# Human-rated similarity on scale of 1-10

word1	word2	similarity
vanish	disappear	9.8
behave	obey	7.3
belief	impression	5.95
muscle	bone	3.65
modest	flexible	0.98
hole	agreement	0.3

# Why not use a thesaurus?

- We don't have a thesaurus for every language
- We can't have a thesaurus for every year



# Distributional Semantics

- What is *tesgüino*?

- (a) *A bottle of tesgüino is on the table*
- (b) *People like tesgüino.*
- (c) *Don't have tesgüino before you drive.*
- (d) *Tesgüino is made out of corn*

- Intuition for algorithm: **Two words are similar if they have similar word contexts.**

- **How do we model context of a word?**

# Two kinds of vector models

- *Sparse representation*

weighted word co-occurrence matrices

- *Dense representation*

Singular Value Decomposition (LSA)

Prediction-based models (Word2vec, GloVe)

Brown clusters

# Two kinds of vector models

- *Sparse representation*

weighted word co-occurrence matrices

# Co-occurrence matrices

- Represent how often a word occurs with other words

**Term-term matrix or word-word co-occurrence matrix or word-context matrix**

# Word-word matrix (or "term-context matrix")

- Fix a context
  - Paragraph
  - Context window of size k


sugar, a sliced lemon, a tablespoonful of  
their enjoyment. Cautiously she sampled her first  
well suited to programming on the digital  
for the purpose of gathering data and

**apricot**  
**pineapple**  
**computer.**  
**information**

jam, a pinch each of,  
and another fruit whose taste she likened  
In finding the optimal R-stage policy from  
necessary for the study authorized in the

	aardvark	computer	data	pinch	result	sugar	...
apricot	0	0	0	1	0	1	
pineapple	0	0	0	1	0	1	
digital	0	2	1	0	1	0	
information	0	1	6	0	4	0	

	aardvark	computer	data	pinch	result	sugar	...
apricot	0	0	0	1	0	1	
pineapple	0	0	0	1	0	1	
digital	0	2	1	0	1	0	
information	0	1	6	0	4	0	



# More common: word-word matrix (or "term-context matrix")

- Each vector of dimension  $|V|$
- The word-word matrix is  $|V| \times |V|$

	Aardvark	Computer	Data	Pie	Result	sugar
pineapple	0	2	8	442	9	25
strawberry	0	0	0	1	60	19
digital	0	2	1683	5	85	4
information	0	3325	3982	5	378	13

Co-occurrence vectors for four words in the Wikipedia corpus, showing six of the dimensions

# More common: word-word matrix (or "term-context matrix")

- We show only 4 x 6, but real matrix is 50,000 x 50,000
- Very sparse (most values are zero)
- That's OK, since there are lots of efficient algorithms for sparse matrices

	Aardvark	Computer	Data	Pie	Result	sugar
cherry	0	2	8	442	9	25
strawberry	0	0	0	1	60	19
digital	0	2	1683	5	85	4
information	0	3325	3982	5	378	13

Co-occurrence vectors for four words in the Wikipedia corpus, showing six of the dimensions



# Reminders from Linear Algebra

$$\text{dot-product}(\vec{v}, \vec{w}) = \vec{v} \cdot \vec{w} = \sum_{i=1}^N v_i w_i = v_1 w_1 + v_2 w_2 + \dots + v_N w_N$$

$$\text{vector length } |\vec{v}| = \sqrt{\sum_{i=1}^N v_i^2}$$

# Measuring similarity

- Given word vectors  $v$  and  $w$
- Need to measure their similarity

We know from linear algebra:

Dot-product (inner product) can measure vector similarity

- High when two vectors have large values in same dimensions
- Low (infact 0) for **orthogonal vectors** with zeros in complementary distribution

# Problem with dot product

$$\text{dot-product}(\vec{v}, \vec{w}) = \vec{v} \cdot \vec{w} = \sum_{i=1}^N v_i w_i = v_1 w_1 + v_2 w_2 + \dots + v_N w_N$$

Vector length:

$$|\vec{v}| = \sqrt{\sum_{i=1}^N v_i^2}$$

$v_i$  is the count for word  $v$  in context  $i$   
 $w_i$  is the count for word  $w$  in context  $i$ .

Vectors are longer if they have higher values in each dimension

- Frequent words have higher values in each dimension
- Dot-product then favors frequent words
- That's bad: We don't want similarity between words sensitive to word frequency (**How do we handle this?**)

# Solution: Cosine Similarity

Divide dot-product by vector lengths!

This turns out to be the cosine of the angle between the vectors.


$$\text{cosine}(\vec{v}, \vec{w}) = \frac{\vec{v} \cdot \vec{w}}{|\vec{v}| |\vec{w}|} = \frac{\sum_{i=1}^N v_i w_i}{\sqrt{\sum_{i=1}^N v_i^2} \sqrt{\sum_{i=1}^N w_i^2}}$$

$$\vec{a} \cdot \vec{b} = |\vec{a}| |\vec{b}| \cos \theta$$

$$\frac{\vec{a} \cdot \vec{b}}{|\vec{a}| |\vec{b}|} = \cos \theta$$

# Cosine as a similarity metric

- -1: vectors point in opposite directions
  - +1: vectors point in same directions
  - 0: vectors are orthogonal
- 
- Frequency is non-negative, so cosine range 0-1



$$\cos(\vec{v}, \vec{w}) = \frac{\vec{v} \cdot \vec{w}}{|\vec{v}| |\vec{w}|} = \frac{\vec{v}}{|\vec{v}|} \cdot \frac{\vec{w}}{|\vec{w}|} = \frac{\sum_{i=1}^N v_i w_i}{\sqrt{\sum_{i=1}^N v_i^2} \sqrt{\sum_{i=1}^N w_i^2}}$$

	large	data	computer
apricot	1	0	0
digital	0	1	2
information	1	6	1

Which pair of words is more similar?

cosine(apricot,information) =

cosine(digital,information) =

cosine(apricot,digital) =

# But raw frequency is a bad representation

- Frequency is clearly useful; if *sugar* appears a lot near *apricot*, that's useful information.
- But overly frequent words like *the*, *it*, or *they* are not very informative about the context

Need a function that resolves this frequency paradox!

# tf-idf: combine two factors

- **tf: term frequency.** frequency count (usually log-transformed):

$$\text{tf}_{t,d} = \begin{cases} 1 + \log_{10} \text{count}(t,d) & \text{if } \text{count}(t,d) > 0 \\ 0 & \text{otherwise} \end{cases}$$

- **idf: inverse document frequency:**

$$\text{idf}_i = \log \left( \frac{N}{\text{df}_i} \right)$$

Total # of docs in collection

# of docs that have word i

Words like "the" or "good" have very low idf

tf-idf value for word  $t$  in document  $d$ :

$$w_{t,d} = \text{tf}_{t,d} \times \text{idf}_t$$



# So far

- Co-occurrence matrix
- Similarity between vectors

# More common: word-word matrix (or "term-context matrix")

- We show only 4 x 6, but real matrix is 50,000 x 50,000
- Very sparse (most values are zero)
- That's OK, since there are lots of efficient algorithms for sparse matrices

	Aardvark	Computer	Data	Pie	Result	sugar
cherry	0	2	8	442	9	25
strawberry	0	0	0	1	60	19
digital	0	2	1683	5	85	4
information	0	3325	3982	5	378	13

Co-occurrence vectors for four words in the Wikipedia corpus, showing six of the dimensions

# Dense Representation

weighted vectors are

- **long** (length  $|V| = 20,000$  to  $50,000$ )
- **sparse** (most elements are zero)

Alternative: learn vectors which are

- **short** (length 200-1000)
- **dense** (most elements are non-zero)

# Sparse vs. Dense vectors

- Why dense vectors?
  - Short vectors may be easier to use as features in machine learning (less weights to tune)
  - Dense vectors may generalize better than storing explicit counts
  - They may do better at capturing synonymy:
    - *car* and *automobile* are synonyms; but are represented as distinct dimensions; this fails to capture similarity between a word with *car* as a neighbor and a word with *automobile* as a neighbor

# Three Methods for Dense Vectors

- Singular Value Decomposition (SVD)
  - A special case of this is called LSA – Latent Semantic Analysis
- “Neural Language Model”-inspired predictive models
  - skip-grams and CBOW
- Brown clustering

# Prediction-based models

- **Word2vec** (Mikolov et al.)
- <https://code.google.com/archive/p/word2vec/>
- **Fasttext** <http://www.fasttext.cc/>
- **Glove** (Pennington, Socher, Manning)
- <http://nlp.stanford.edu/projects/glove/>

# How to learn word2vec (skip-gram) embeddings


- Start with  $V$  random 300-dimensional vectors as initial embeddings
- Use logistic regression (the second most basic classifier used in machine learning after naïve Bayes)
  - Take a corpus and take pairs of words that co-occur as positive examples
  - Take pairs of words that don't co-occur as negative examples
  - Train the classifier to distinguish these by slowly adjusting all the embeddings to improve the classifier performance
  - Throw away the classifier code and keep the embeddings.

- Instead of **counting** how often each word  $w$  occurs near "*apricot*"
- Train a classifier on a binary **prediction** task:
  - Is  $w$  likely to show up near "*apricot*"?
- We don't actually care about this task
  - But we'll take the learned classifier weights as the word embeddings



# How to learn word2vec (skip-gram)

## Skip-gram



<b><i>word</i></b>	<b>dim0</b>	<b>dim1</b>	<b>dim2</b>	<b>dim3</b>	<b>...</b>	<b>dim300</b>
<i>today</i>	0.35	-1.3	2.2	0.003		
<i>cat</i>	-3.1	-1.7	1.1	-0.56		
<i>sleep</i>	0.55	3.0	2.4	-1.2		
<i>watch</i>	-0.09	0.8	-1.8	2.9		
<i>bird</i>	2.0	0.16	-1.9	2.3		

# Evaluating embeddings

- Have a vocabulary and a set of word embeddings
- Need to know how good the vectors are
- Intrinsic:
  - Compare to human scores on word similarity-type tasks:
    - WordSim-353 dataset
      - environment – ecology      8.8
- Extrinsic

# Properties of embeddings

Similarity depends on window size  $C$

- $C = \pm 2$  The nearest words to *Hogwarts*:
  - *Sunnydale*
  - *Evernight*
- $C = \pm 5$  The nearest words to *Hogwarts*:
  - *Dumbledore*
  - *Malfoy*
  - *halfblood*

# Properties of embeddings



word = “sweden”

Word	Cosine distance
norway	0.760124
denmark	0.715460
finland	0.620022
switzerland	0.588132
belgium	0.585835
netherlands	0.574631
iceland	0.562368
estonia	0.547621
slovenia	0.531408

# Analogy: Embeddings capture relational meaning

$\text{vector}(\text{'king'}) - \text{vector}(\text{'man'}) + \text{vector}(\text{'woman'}) \approx \text{vector}(\text{'queen'})$

$\text{vector}(\text{'Paris'}) - \text{vector}(\text{'France'}) + \text{vector}(\text{'Italy'}) \approx \text{vector}(\text{'Rome'})$





**Figure 6.1** A two-dimensional (t-SNE) projection of embeddings for some words and phrases, showing that words with similar meanings are nearby in space. The original 60-dimensional embeddings were trained for sentiment analysis. Simplified from [Li et al. \(2015\)](#).

# Summary

- Distributional (vector) Models of meaning
  - Sparse vectors:
    - tf-idf weighted word-word co-occurrence matrix
  - Dense vectors:
    - Prediction-based: Word2vec